



# MECH5170M Connected and Autonomous Vehicles Systems

Machine learning algorithms

Kris Kubiak ( [k.kubiak@leeds.ac.uk](mailto:k.kubiak@leeds.ac.uk) )

## ARTIFICIAL INTELLIGENCE VS MACHINE LEARNING VS DEEP LEARNING

### ① Artificial Intelligence

Development of smart systems and machines that can carry out tasks that typically require human intelligence

### ② Machine Learning

Creates algorithms that can learn from data and make decisions based on patterns observed  
Require human intervention when decision is incorrect

### ③ Deep Learning

Uses an artificial neural network to reach accurate conclusions without human intervention



3

UNIVERSITY OF LEEDS

# Machine Learning

# Research Fields

Statistics / Statistical Learning

Data Mining

Pattern Recognition

Artificial Intelligence

Databases

Signal Processing



# Machine Learning Algorithm Cheat Sheet

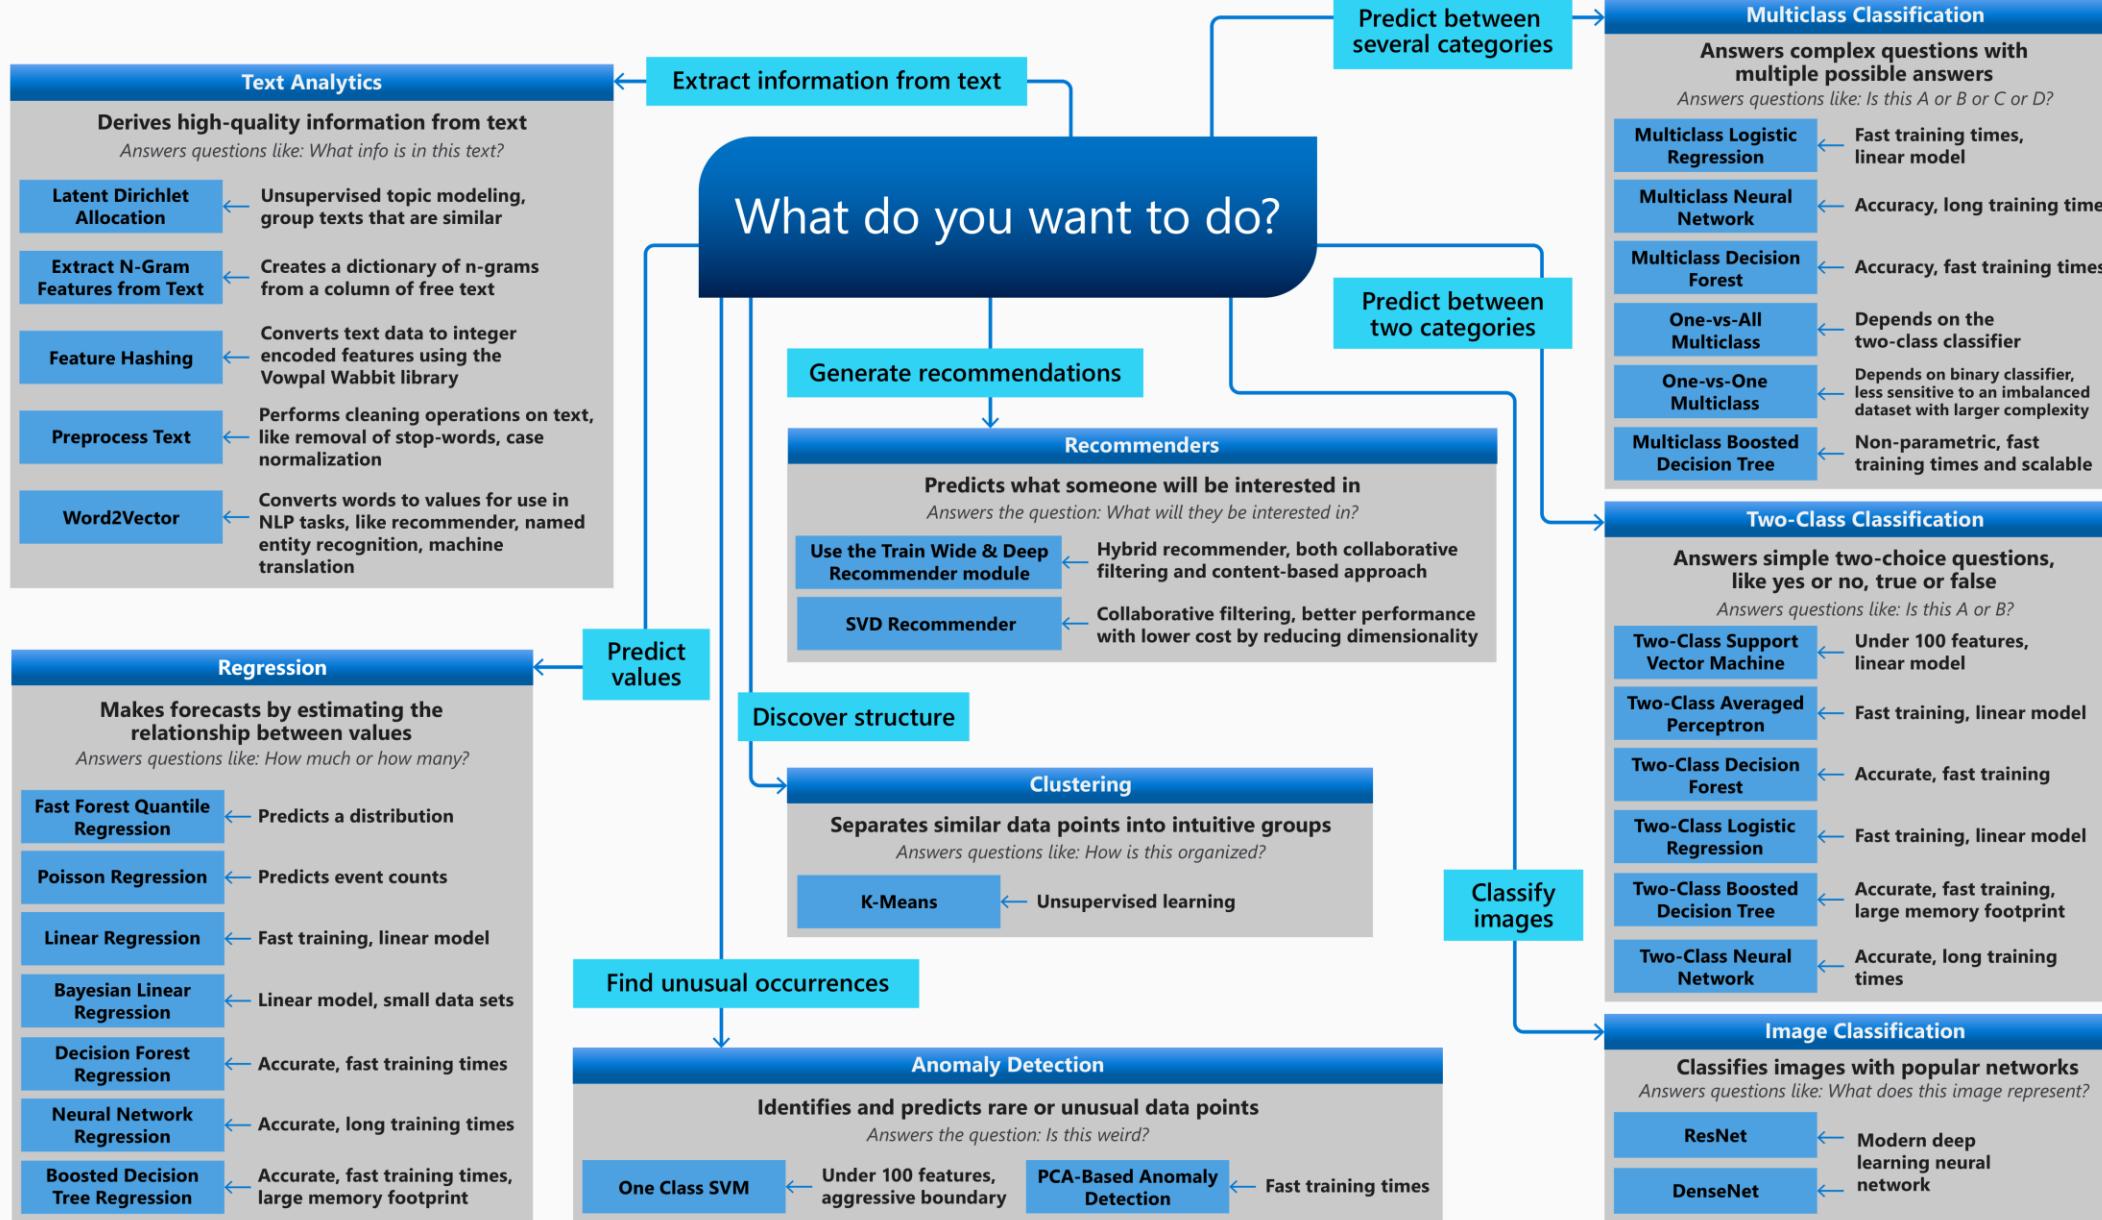
This cheat sheet helps you choose the best machine learning algorithm for your predictive analytics solution. Your decision is driven by both the nature of your data and the goal you want to achieve with your data.

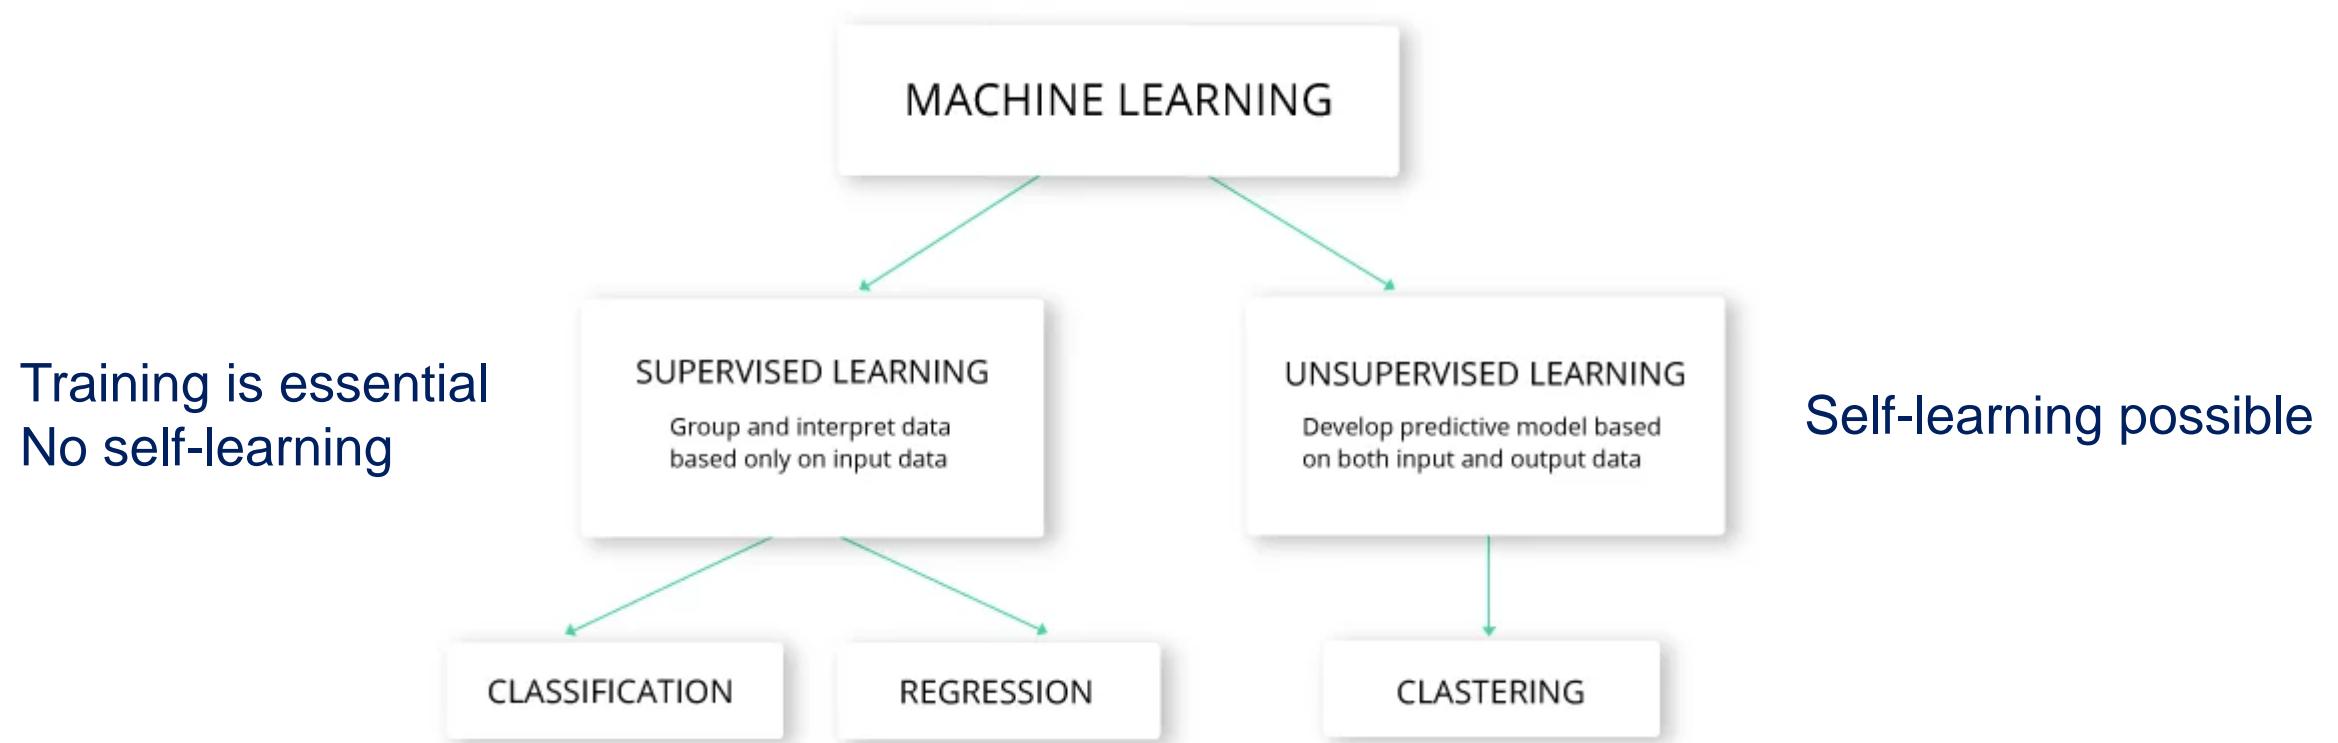


Microsoft Azure



5





Object Detection

Object Classification

Distance, Speed, Direction Estimation

Path Finding

Trajectory planning

Motor/Battery Monitoring

Anomaly Detection

Clustering

Multiclass Classification

Regression - value prediction

Optimisation Algorithms

Collision avoidance (prediction)

Regression

Pattern Recognition

# Model Building Process

## Data Splitting

- Allocate data to different tasks
  - Model training
  - Performance evaluation
- Define Training, Validation and Test sets

## Feature Selection (Review the decision made previously)

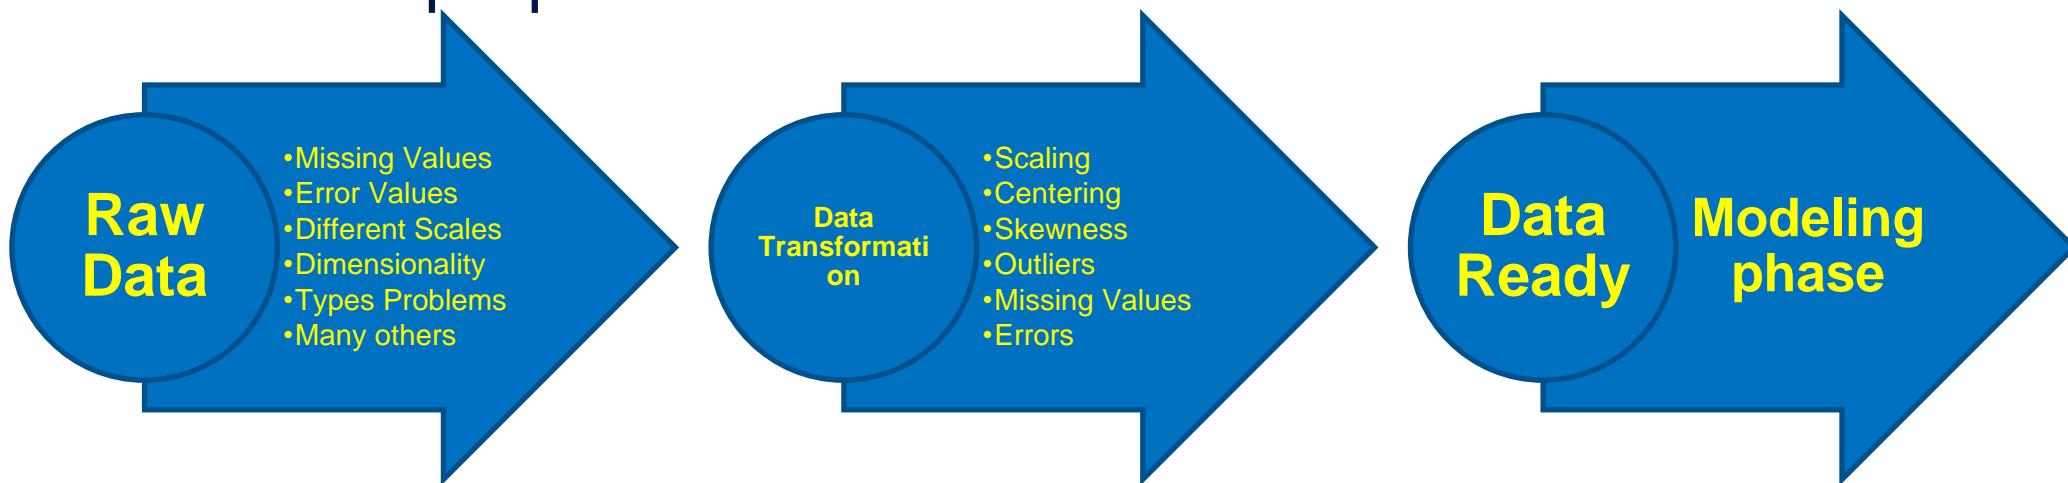
## Estimating Performance

- Visualization of results – discover interesting areas of the problem space
- Statistics and performance measures

## Evaluation and Model selection

- No prior assumptions should be made
- Avoid use of favorite models, evaluate and select the best performing one

- Needed for several reasons
  - Some Models have strict data requirements
    - Scale of the data, data point intervals, etc.
  - Some characteristics of the data may impact dramatically on the model performance
- Time on data preparation should not be underestimated



## Supervised Learning

- For every example in the data there is always a predefined outcome
- Models the relations between a set of descriptive features and a target (Fits data to a function)
- 2 groups of problems:
  - Classification
  - Regression

# Supervised Learning

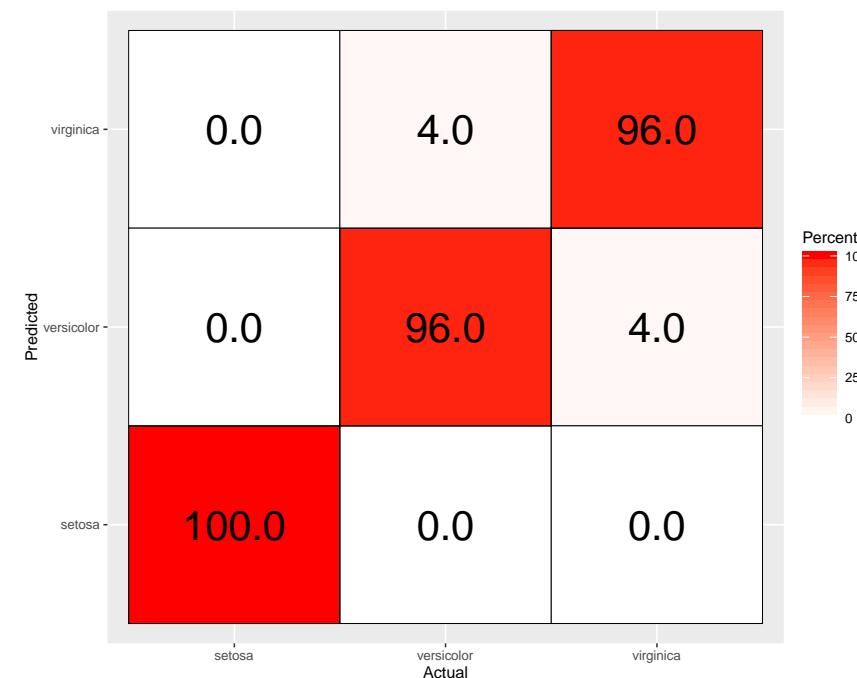
## Classification

- Predicts which class a given sample of data (sample of descriptive features) is part of (**discrete value**)

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
4	4.6	3.1	1.5	0.2	setosa
5	5.0	3.6	1.4	0.2	setosa
6	5.4	3.9	1.7	0.4	setosa
7	4.6	3.4	1.4	0.3	setosa
8	5.0	3.4	1.5	0.2	setosa
9	4.4	2.9	1.4	0.2	setosa
10	4.9	3.1	1.5	0.1	setosa

## Regression

- Predicts continuous values

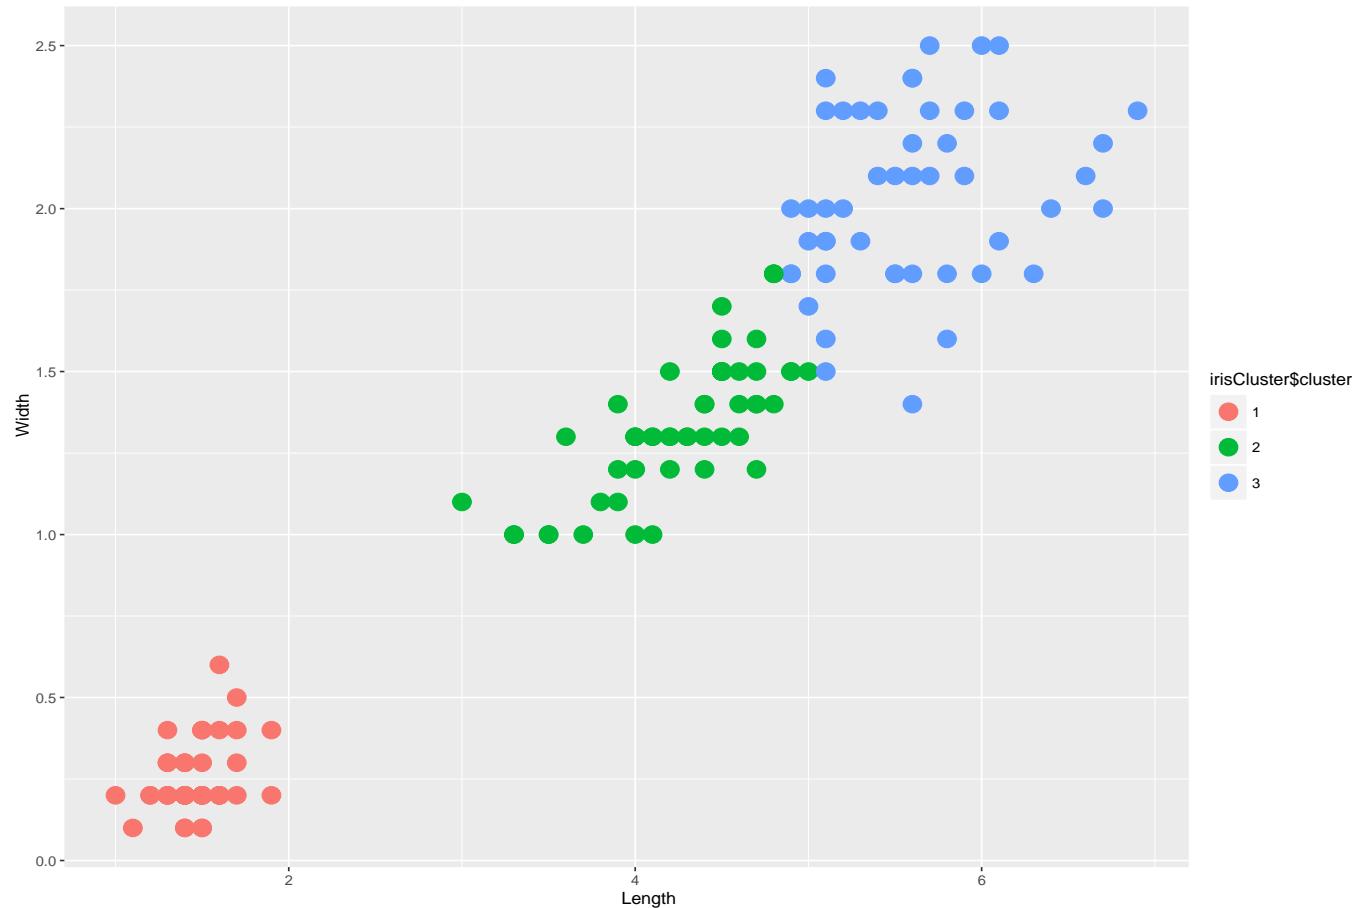


# Unsupervised Learning

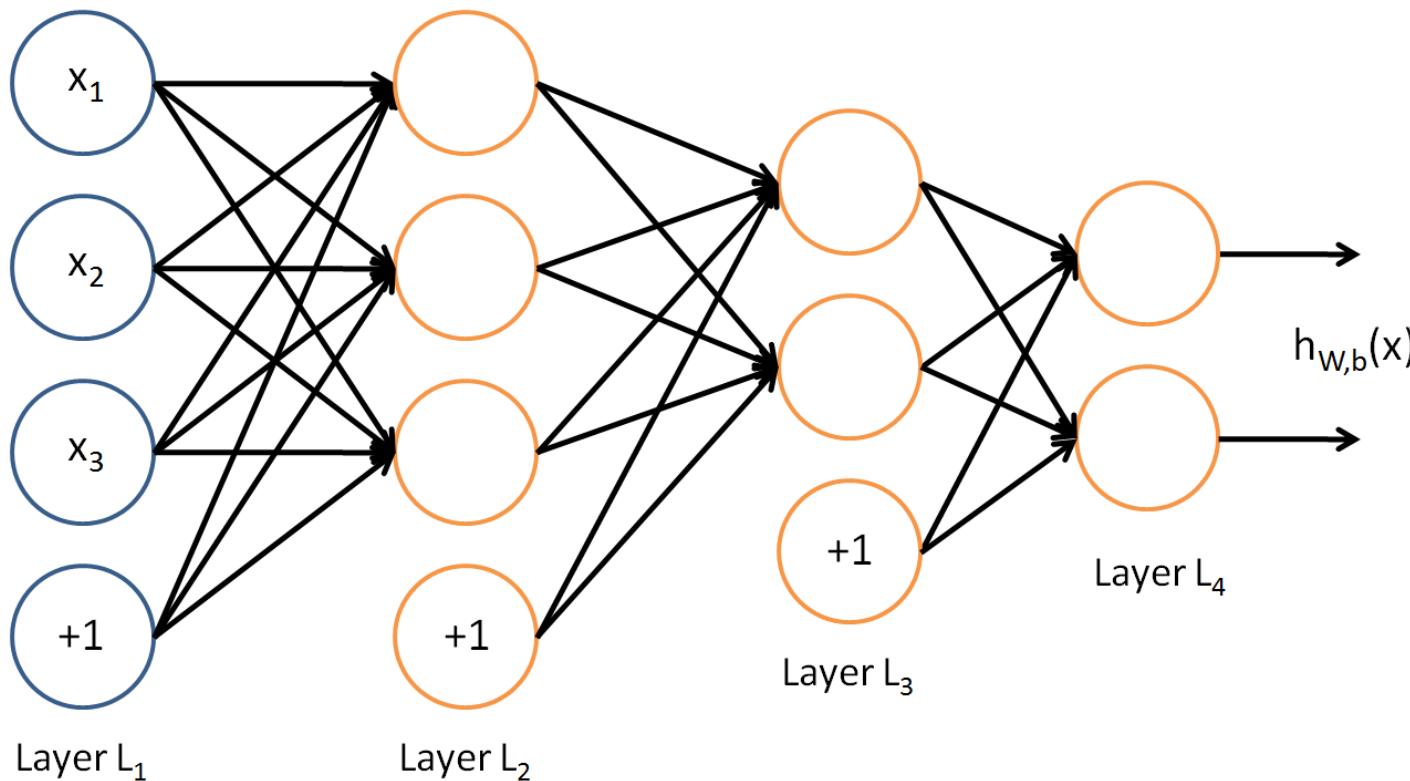
## Unsupervised Learning

- There are not predefined and known set of outcomes
- Look for hidden patterns and relations in the data
- A typical example: Clustering

	Length	Width	Length	Width
1	5.1	3.5	1.4	0.2
2	4.9	3.0	1.4	0.2
3	4.7	3.2	1.3	0.2
4	4.6	3.1	1.5	0.2
5	5.0	3.6	1.4	0.2
6	5.4	3.9	1.7	0.4
7	4.6	3.4	1.4	0.3
8	5.0	3.4	1.5	0.2
9	4.4	2.9	1.4	0.2
10	4.9	3.1	1.5	0.1



# Multilayer network



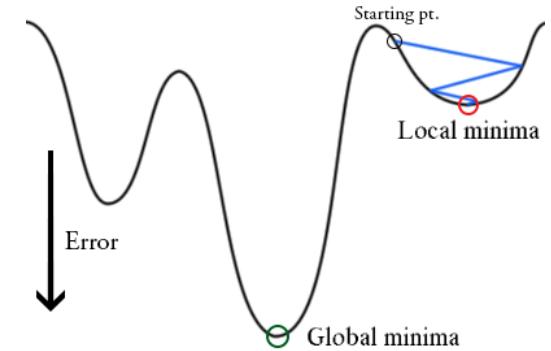
How to determine weights?

# Training: backpropagation

Initialize weights "randomly"

For all training epochs

- for all input-output in training set
  - using input, compute output (forward)
  - compare computed output with training output
  - adapt weights (backward) to improve output
- if accuracy is good enough, stop



Learning from data is used in situations where we don't have any analytical solution, but we do have data that we can use to construct an empirical solution

The basic premise of learning from data is the use of a set of observations to uncover an underlying process.

# The Learning Problem

Suppose we observe the output space  $Y_i$  and the input space

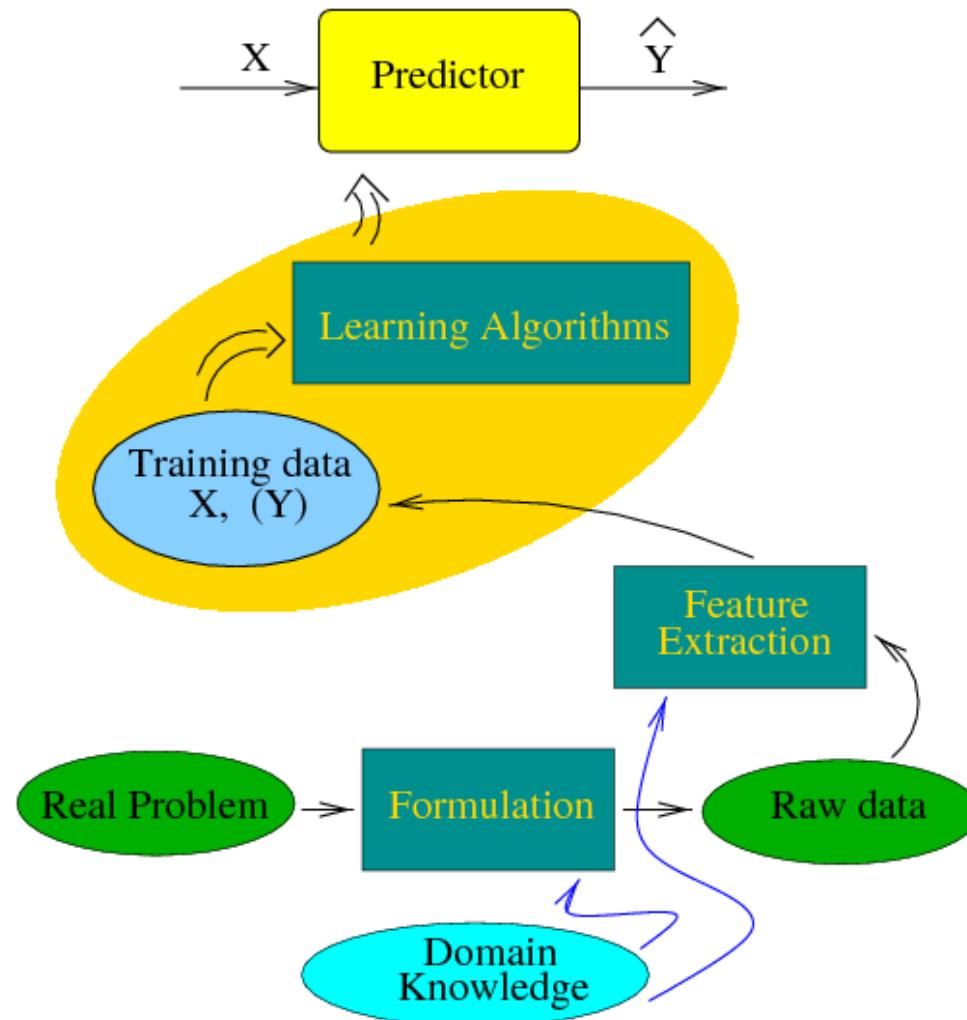
$$X_i = (X_{i1}, \dots, X_{ip}) \quad \forall i = 1, \dots, n$$

We believe that there is a *relationship* between  $Y$  and at least one of the  $X$ 's.

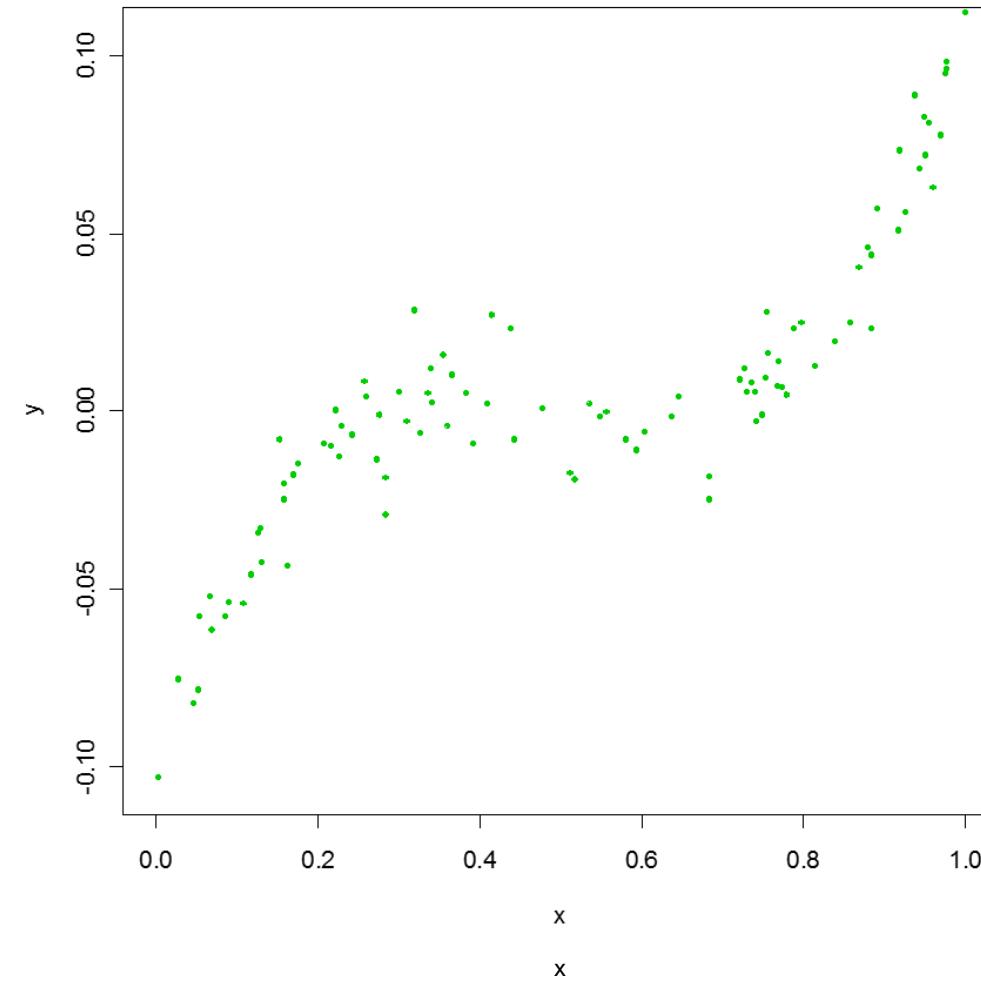
We can model the relationship as:

$$Y_i = f(X_i) + \varepsilon_i$$

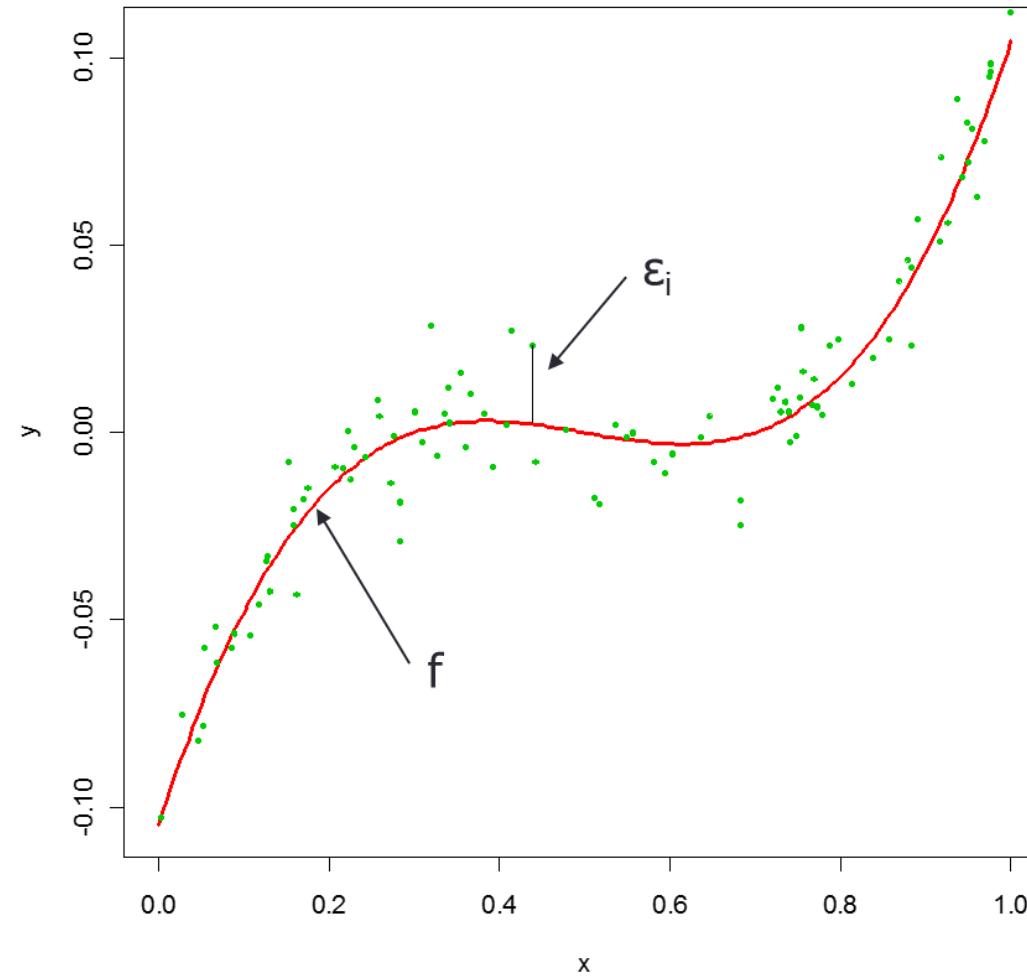
where  $f$  is an unknown function and  $\varepsilon$  is a random error (noise) term, independent of  $X$  with mean zero.



# The Learning Problem: Example

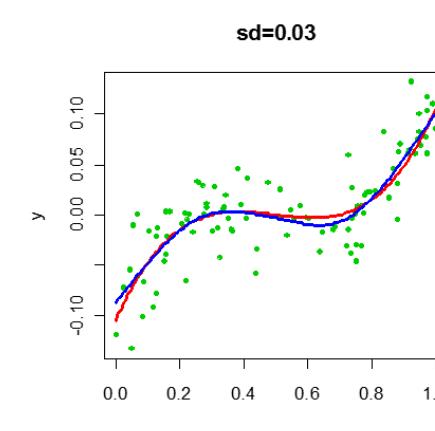
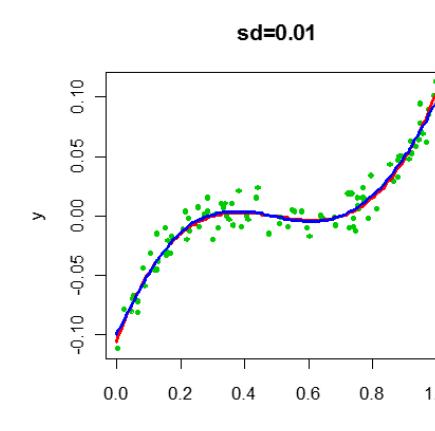
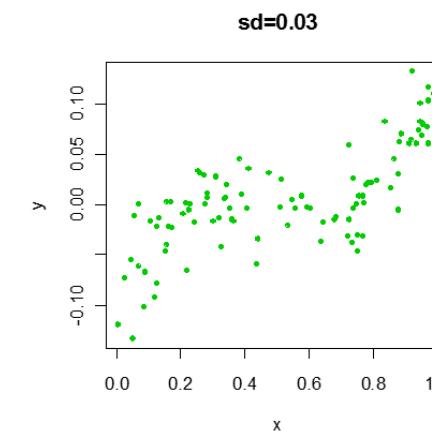
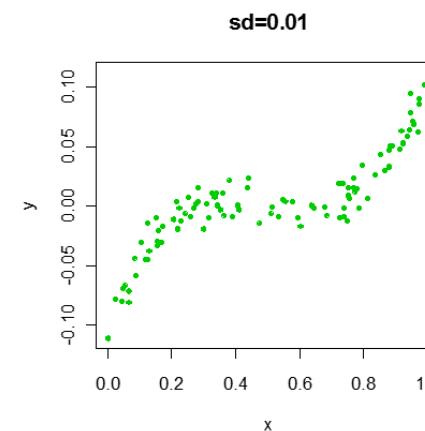
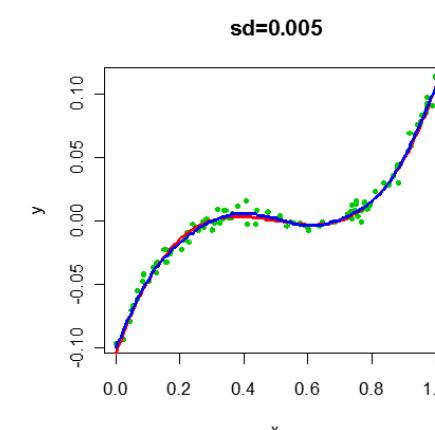
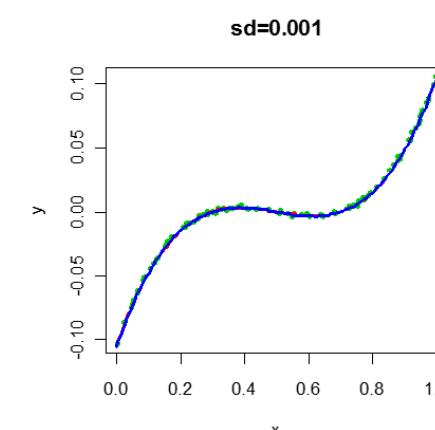
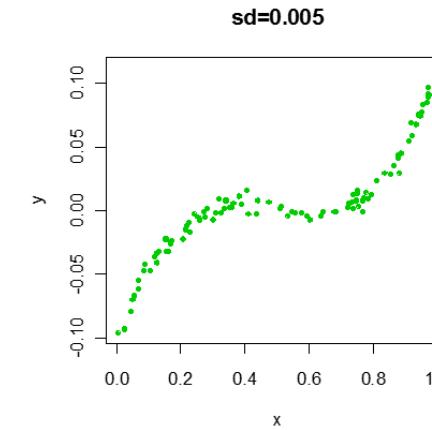
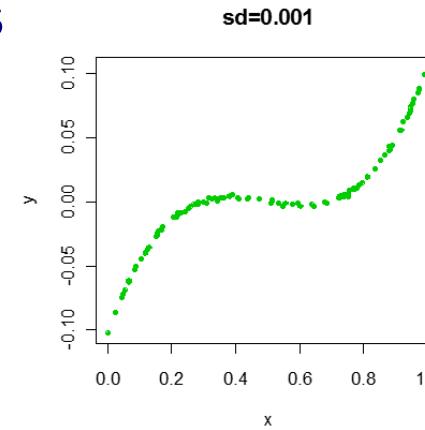


# The Learning Problem: Example



# The Learning Problem: Example

Different estimates for the target function  $f$  that depend on the standard deviation of the  $\varepsilon$ 's



## Why do we estimate $f$ ?

We use modern machine learning methods to estimate  $f$  by *learning* from the data.

The target function  $f$  is unknown.

We estimate  $f$  for two key purposes:

- Prediction
- Inference

# Prediction

By producing a good estimate for  $f$  where the variance of  $\varepsilon$  is not too large, then we can make accurate predictions for the response variable,  $Y$ , based on a new value of  $X$ .

We can predict  $Y$  using  $\hat{Y} = \hat{f}(X)$

where  $\hat{f}$  represents our estimate for  $f$ , and  $\hat{Y}$  represents the resulting prediction for  $Y$ .

The accuracy of  $\hat{Y}$  as a prediction for  $Y$  depends on:

- Reducible error
- Irreducible error

Note that  $\hat{f}$  will not be a perfect estimate for  $f$ ; this inaccuracy introduces error.

This error is **reducible** because we can potentially improve the accuracy of the estimated (i.e. hypothesis) function  $\hat{f}$  by using the most appropriate learning technique to estimate the target function  $f$ .

Even if we could perfectly estimate  $f$ , there is still **variability associated with  $\epsilon$**  that affects the accuracy of predictions = **irreducible** error.

# Prediction

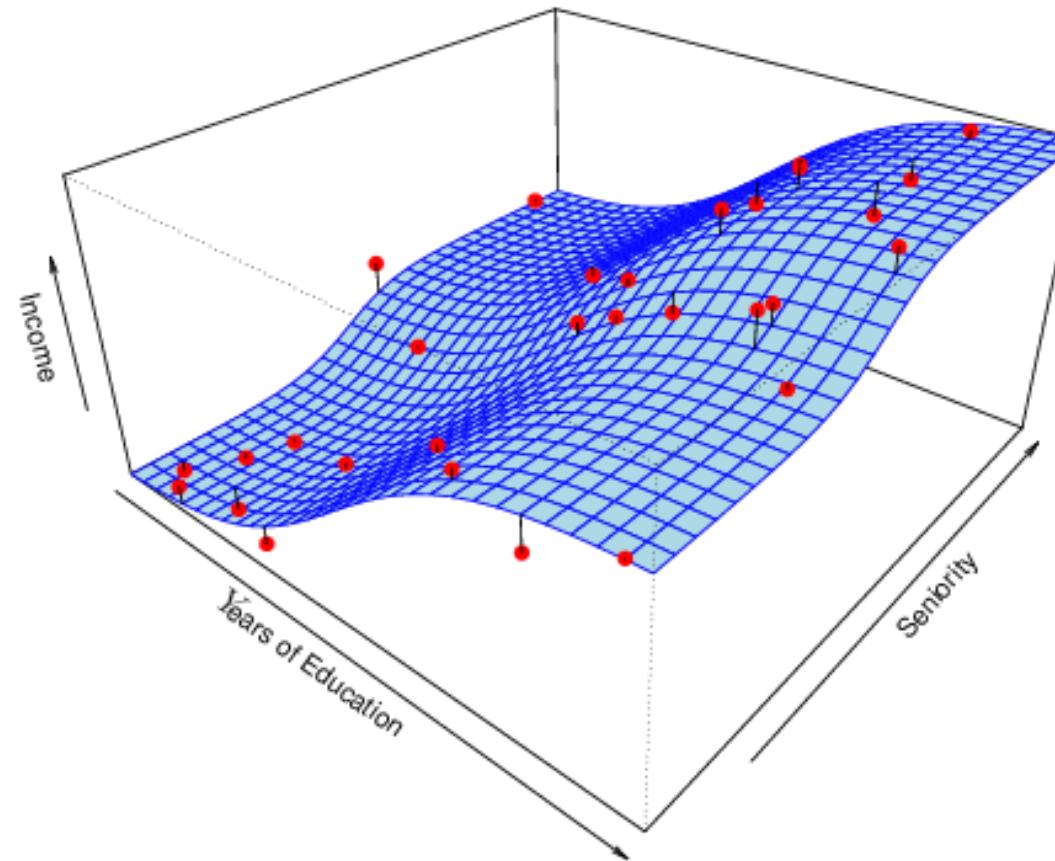
Average of the squared difference between the predicted and actual value of  $Y$ .

$\text{Var}(\epsilon)$  represents the *variance* associated with  $\epsilon$ .

$$E[(Y - \hat{f}(X))^2 | X = x] = \underbrace{[f(x) - \hat{f}(x)]^2}_{\text{Reducible}} + \underbrace{\text{Var}(\epsilon)}_{\text{Irreducible}}$$

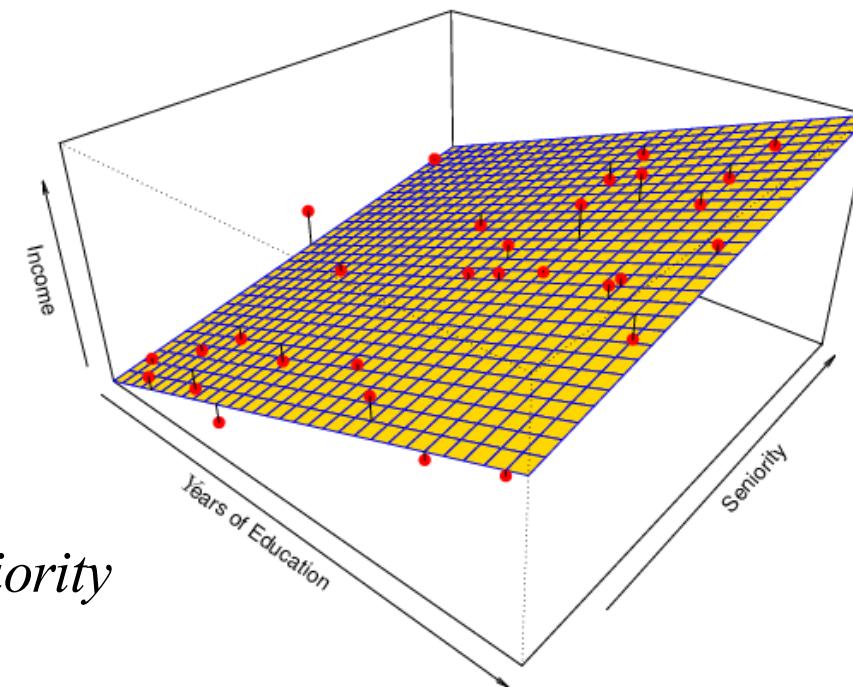
Our aim is to minimize the reducible error!!

# Example: Income vs. Education Seniority



## Example: OLS Regression Estimate

Even if the standard deviation is low, we will still get a bad answer if we use the incorrect model.



$$f = b_0 + b_1 \cdot \text{Education} + b_2 \cdot \text{Seniority}$$

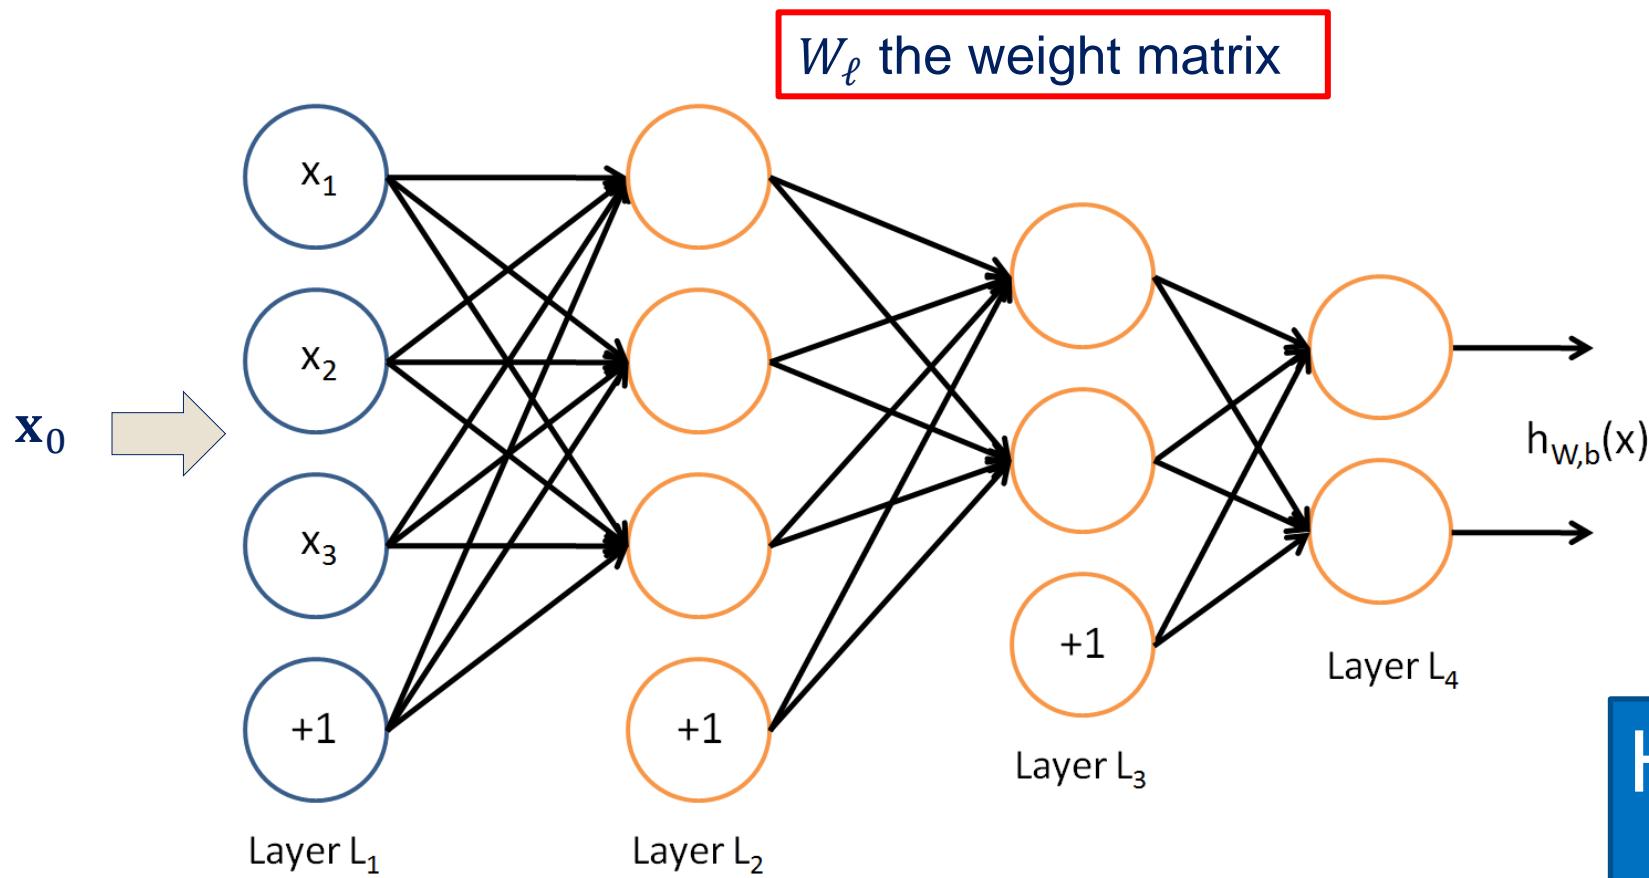
# Machine Learning Conclusions

- Machine learning
  - Supervised (Training before deployment)
  - Unsupervised (Self-learning capabilities)
- Clustering
- Classification
- Regression, Prediction



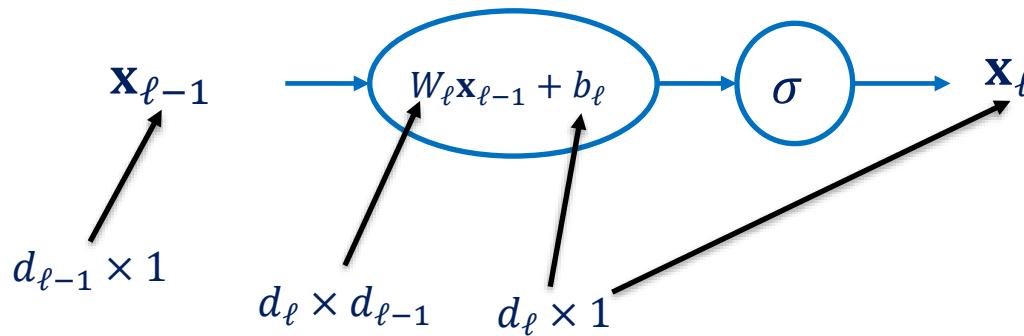
# Neural Networks

# Multilayer network



How to determine  
weights?

# Basics of Neural networks



**Compact representation of the  $\ell^{th}$  layer of a neural network**

A feedforward neural network with  $m$  hidden layers is defined as follows:

- $\mathbf{x}_0$  = Input to the network
- $\mathbf{x}_{m+1}$  = Output of the network
- $\mathbf{x}_\ell$  = Output of the  $\ell^{th}$  layer
- And,  $\mathbf{x}_\ell = f_\ell(\mathbf{x}_{\ell-1})$
- Where,  $f_\ell = \sigma_\ell(W_\ell \mathbf{x}_{\ell-1} + b_\ell)$
- $\sigma_\ell$  is known as the activation function,  $W_\ell$  the weight matrix and,  $b_\ell$  the bias term of  $\ell^{th}$  layer

# Basics of Training a NN

Compute “how off” you are from the given example

$$J(W, b; x, y) = \frac{1}{2} \|h_{W,b}(x) - y\|^2$$

$J$  is the cost function for a given set of weights  $W$ , biases  $b$  and a training example  $x$  and  $y$

Add a regularisation term (to avoid overfitting). This is a term proportional to the sum of squares of all elements in the weight matrices

Use gradient descent to minimise cost

- update each “variable” (weights/biases) :  $W_{i,j} = W_{i,j} - \alpha \frac{\partial}{\partial W_{i,j}} J(W, b)$
- Here,  $\alpha$  is a learning rate, and the partial derivative is computed by back-propagation (chain rule applied to the function represented by the NN)



Inspired by visual cortex in animals

Learns image filters that were previously hand-engineered

Basic intuitions for CNNs:

- Images are too large to be monolithically processed by a feedforward neural network ( $1000 \times 1000$  image =  $10^6$  inputs, which means the weight matrix for the second layer is proportional to at least  $10^{12}$ )
- Data in an image is spatially correlated

CNN divided into several layers with different purposes

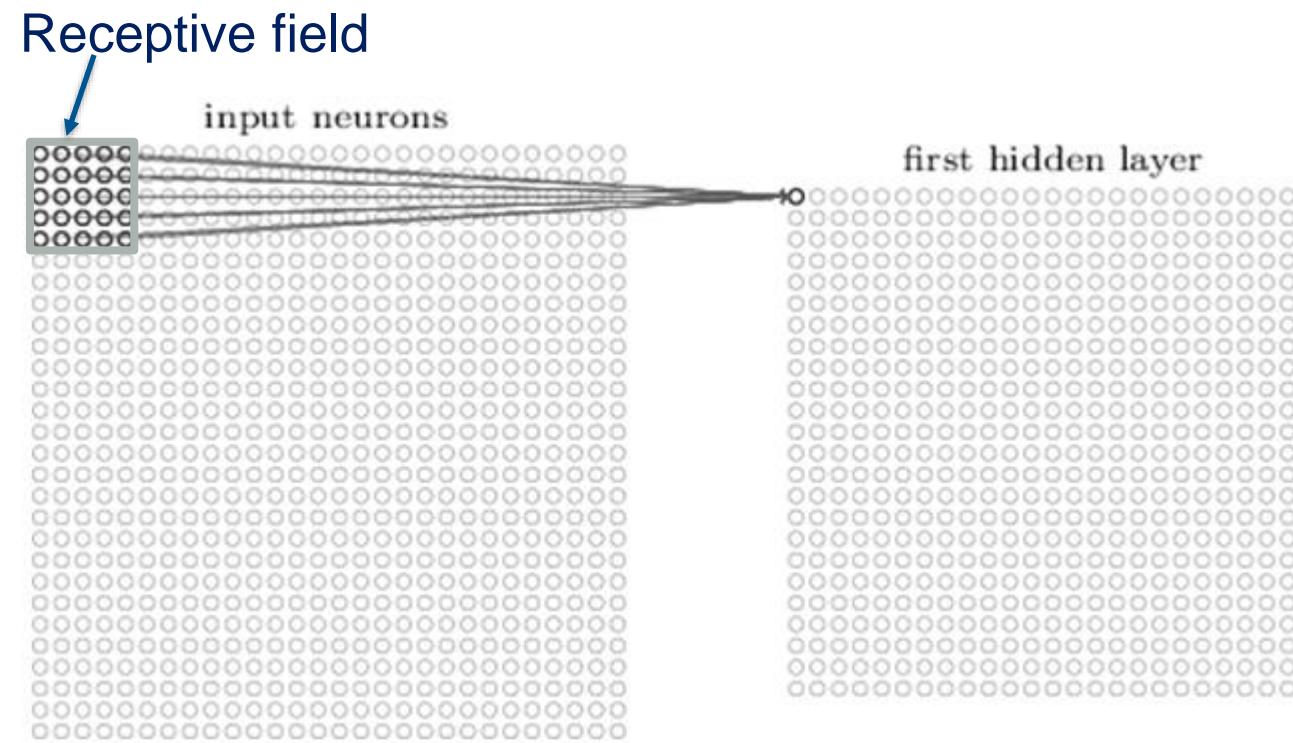
# Convolutional layer

First layer is a convolutional layer

Convolutional layer contains neurons associated with sub-regions of the original image

Each sub-region is called receptive field

Convolves the weights of the convolutional layer with each cell in receptive field to obtain *activation map* or *feature map*



# “Convolution”

**Convolution** of a 2-D image  $I$  with a 2-D kernel  $K$  defined as:

$$S(i, j) = (I * K)(i, j) = \sum_m \sum_n I(m, n)K(i - m, j - n)$$

Most neural network libraries do not use convolution, but instead implement **cross-correlation**, i.e.

$$S(i, j) = \sum_m \sum_n I(i + m, j + n)K(m, n)$$

The kernel function  $K$  usually defines the receptive field

# Purpose of convolutional layer

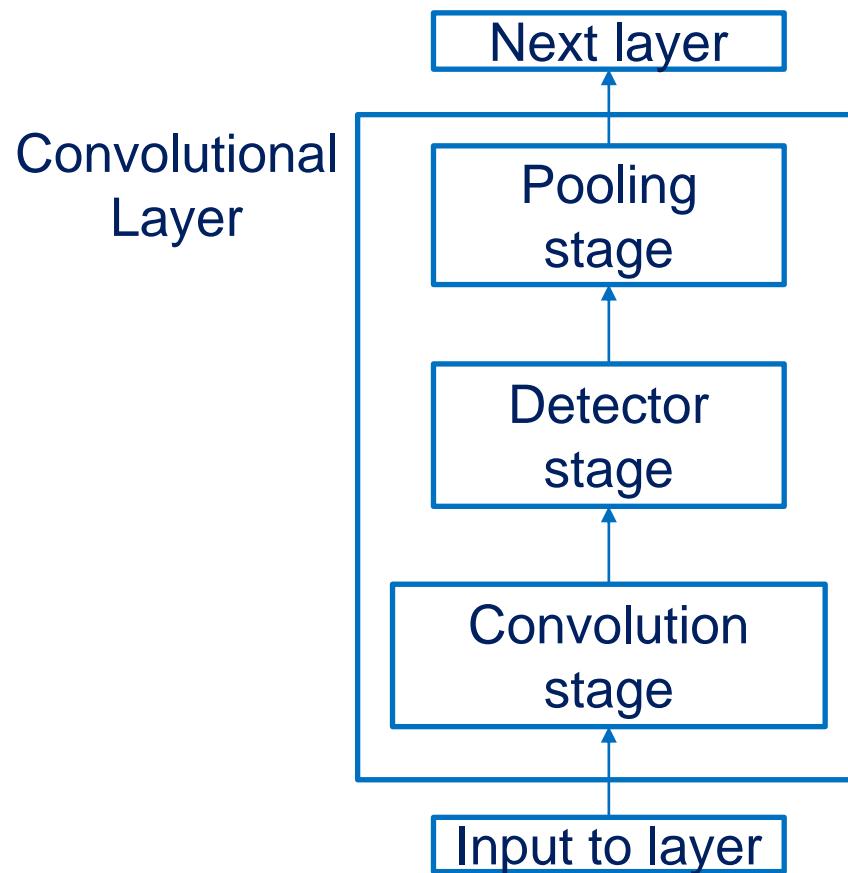
Convolutional layer **applies a *filter* to the pixels** within its **receptive field**

This allows **identifying low-level features** (curves, straight lines, etc.)

The outputs of the first convolutional layer can be thought of as having a “**high value**” when a particular **feature is detected** and a “**low value**” otherwise

**Second convolutional layer** allows learning higher-level features (e.g. semi-circles, angles etc.)

- Second convolutional layer has a bigger receptive field (as it is able to simultaneously correlate over outputs of first layer)
- By “convolving” over the feature map, **output of second layer tries to connect higher level features**



## CNN architecture

- **Small number of layers**, where each layer contains stages
- **Convolution stage** performs convolution
- **Detector stage** uses a nonlinearity such as a rectified linear unit, i.e.  $\max(x, 0)$
- **Pooling stage** performs a suitable pooling operation

## Pooling stage

Pooling function **replaces the output** of a layer at a certain location **with a summary statistic** of the nearby outputs

For example: **max pooling** reports **maximum output** within a rectangular neighbourhood

Other pooling functions include **average**,  **$L^2$  norm** (Euclidean distance), **weighted average** etc.

Pooling helps representation approximately **invariant to small translations**

By **pooling** over outputs of **different convolutions**, features can learn which transformations to become invariant to (e.g. rotation etc.)

# Fully connected layers

CNNs may have some **fully connected layers before the final output**

These layers allows performing **higher-level reasoning** over different **features learned by the previous convolutional layers**

Various kind of **convolution** functions, **pooling** functions and **detection** functions are possible, giving rise to many different flavors

**Number of convolutional layers** can be varied depending on complexity of features to be learned

# YOLO algorithm (You Only Look Once)

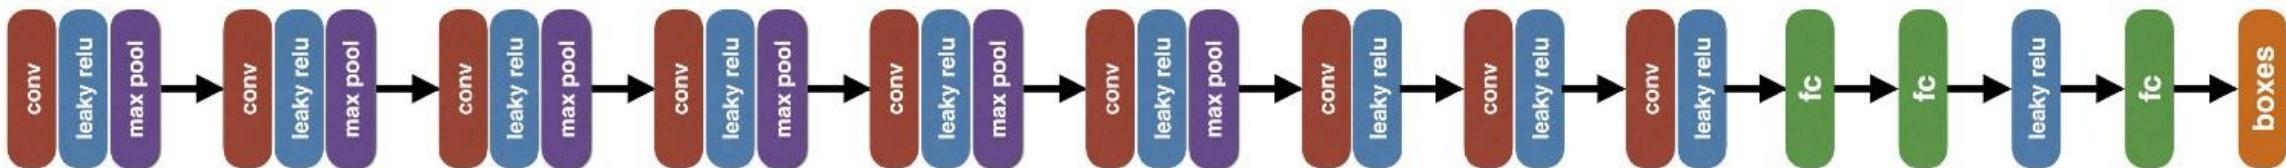
YOLO is one of the fastest real-time detection algorithms

R-CNN etc. are algorithms that leverage classifiers and localisers to perform detection

YOLO applies a single neural network to the entire image

Network divides image into regions and predicts bounding boxes and probabilities for each region

Bounding boxes are weighted by predicted probabilities



# Annotations

- category
- attribute
- visibility
- instance
- sensor
- calibrated\_sensor
- ego\_pose
- log
- scene
- sample
- sample\_data
- sample\_annotation
- map

animal
debris
<b>pushable_pullable</b>
bicycle_rack
ambulance
police
barrier
bicycle
bus.bendy
bus.rigid
<b>car</b>
construction
motorcycle
adult
child
<b>construction_worker</b>
police_officer
<b>personal_mobility</b>
stroller
wheelchair
trafficcone
<b>trailer</b>
<b>truck</b>

<b>type</b>	Car, Van, Truck, Pedestrian, Person_sitting, Cyclist, Tram, Misc or DontCare
truncated	0 to 1, where truncated refers to the object leaving image boundaries
occluded	0 = fully visible, 1 = partly occluded, 2 = largely occluded, 3 = unknown
alpha	Observation angle of object, ranging [-pi..pi]
bbox	2D bounding box of object in the image
dimensions	3D object dimensions: height, width, length
location	3D object location x,y,z in camera coordinate
rotation_y	Rotation ry around Y-axis in camera coordinates [-pi..pi]

# Conclusions

## Machine learning:

- Supervised (Training before deployment)
- Unsupervised (Self-learning capabilities)

## Neural Network - Main steps:

- Object recognition
- Object perception
- Objects annotations

## Image recognition:

- YOLO is often used for the image processing
- Bounding boxes are used, and annotated

# ANY QUESTIONS

???